

# Relational Algebra

Dr Paolo Guagliardo



THE UNIVERSITY of EDINBURGH  
**informatics**

Fall 2020 (v20.2.0)

## Data model

- ▶ A **relation** is a **set** of **records**  
over the same **set** of (distinct) attribute names
- ▶ A **record** is a **total function** from attribute names to values

## Schema

- ▶ Set of **relation names**
- ▶ Set of **distinct attributes** for each table  
**Note that columns are not ordered**

## Instance

- ▶ Actual data (that is, the records in each relation)

# Relational algebra

## Procedural query language

A relational algebra expression

- ▶ takes as input one or more relations
- ▶ applies a **sequence of operations**
- ▶ returns a relation as output

### Operations:

Projection ( $\pi$ )

Selection ( $\sigma$ )

Product ( $\times$ )

Renaming ( $\rho$ )

Union ( $\cup$ )

Intersection ( $\cap$ )

Difference ( $-$ )

The application of each operation results in a new (virtual) relation that can be used as input to other operations

## Projection

- ▶ **Vertical operation**: choose some of the columns
- ▶ **Syntax**:  $\pi_{\text{set of attributes}}(\text{relation})$
- ▶  $\pi_{A_1, \dots, A_n}(R)$  takes only the values of attributes  $A_1, \dots, A_n$  for each tuple in  $R$

Customer

CustID	Name	City	Address
cust1	Renton	Edinburgh	2 Wellington Pl
cust2	Watson	London	221B Baker St
cust3	Holmes	London	221B Baker St

$\pi_{\text{Name, City}}(\text{Customer})$

Name	City
Renton	Edinburgh
Watson	London
Holmes	London

## Selection

- ▶ **Horizontal operation**: choose rows satisfying some condition
- ▶ **Syntax**:  $\sigma_{\text{condition}}(\text{relation})$
- ▶  $\sigma_{\theta}(R)$  takes only the tuples in  $R$  for which  $\theta$  is satisfied

**term** := attribute | constant

$\theta$  := **term** **op** **term** with **op**  $\in \{=, \neq, >, <, \geq, \leq\}$   
|  $\theta \wedge \theta$  |  $\theta \vee \theta$  |  $\neg\theta$

## Example of selection

Customer

CustID	Name	City	Age
cust1	Renton	Edinburgh	24
cust2	Watson	London	32
cust3	Holmes	London	35

$\sigma_{\text{City} \neq \text{'Edinburgh'} \wedge \text{Age} < 33}(\text{Customer})$

CustID	Name	City	Age
cust2	Watson	London	32

## Efficiency (1)

Consecutive selections can be combined into a single one:

$$\sigma_{\theta_1}(\sigma_{\theta_2}(R)) \equiv \sigma_{\theta_1 \wedge \theta_2}(R)$$

↳ is equivalent to: two queries  $Q_1$  and  $Q_2$  are **equivalent** if  $Q_1$  returns the same answers as  $Q_2$  on every database

### Example

$$Q_1 = \sigma_{\text{City} \neq \text{'Edinburgh'}}(\sigma_{\text{Age} < 33}(\text{Customer}))$$

$$Q_2 = \sigma_{\text{City} \neq \text{'Edinburgh'} \wedge \text{Age} < 33}(\text{Customer})$$

$Q_1 \equiv Q_2$  but  $Q_2$  **faster** than  $Q_1$  in general

## Efficiency (2)

Projection can be pulled in front of selection

$$\sigma_{\theta}(\pi_{\alpha}(R)) \equiv \pi_{\alpha}(\sigma_{\theta}(R))$$

**only if all attributes mentioned in  $\theta$  appear in  $\alpha$**

### Example

$$Q_1 = \pi_{\text{Name, City, Age}}(\sigma_{\text{City} \neq \text{'Edinburgh'} \wedge \text{Age} < 33}(\text{Customer}))$$

$$Q_2 = \sigma_{\text{City} \neq \text{'Edinburgh'} \wedge \text{Age} < 33}(\pi_{\text{Name, City, Age}}(\text{Customer}))$$

**Question:** Which one is more efficient?

## Cartesian product

$R \times S$  **concatenates** each tuple of  $R$  with all the tuples of  $S$

**Note:** the relations must have **disjoint** sets of attributes

### Example

$R$	<b>A</b>	<b>B</b>	$\times$	$S$	<b>C</b>	<b>D</b>	$=$	$R \times S$	<b>A</b>	<b>B</b>	<b>C</b>	<b>D</b>
	1	2			1	a			1	2	1	a
	3	4			2	b			1	2	2	b
					3	c			1	2	3	c
									3	4	1	a
									3	4	2	b
									3	4	3	c

### Expensive operation:

- ▶  $\text{card}(R \times S) = \text{card}(R) \times \text{card}(S)$
- ▶  $\text{arity}(R \times S) = \text{arity}(R) + \text{arity}(S)$

## Joining relations

Combining Cartesian product and selection

**Customer:** ID, Name, City, Address

**Account:** Number, Branch, CustID, Balance

We can join customers with the accounts they own as follows

$$\sigma_{\text{ID}=\text{CustID}}(\text{Customer} \times \text{Account})$$

## Renaming

Gives a new name to some of the attributes of a relation

**Syntax:**  $\rho_{\text{replacements}}(\text{relation})$ ,  
where a replacement has the form  $A \rightarrow B$

$$\rho_{A \rightarrow A', C \rightarrow D} \left( \begin{array}{ccc} \mathbf{A} & \mathbf{B} & \mathbf{C} \\ \hline a & b & c \\ 1 & 2 & 3 \end{array} \right) = \begin{array}{ccc} \mathbf{A'} & \mathbf{B} & \mathbf{D} \\ \hline a & b & c \\ 1 & 2 & 3 \end{array}$$

### Example

**Customer:** **CustID**, Name, City, Address

**Account:** Number, Branch, **CustID**, Balance

$$\sigma_{\text{CustID}=\text{CustID}'}(\text{Customer} \times \rho_{\text{CustID} \rightarrow \text{CustID}'}(\text{Account}))$$

## Natural join

Joins two tables on their **common attributes**

### Example

**Customer:** **CustID**, Name, City, Address

**Account:** Number, Branch, **CustID**, Balance

$$\text{Customer} \bowtie \text{Account} \equiv$$

$$\pi_{X \cup Y}(\sigma_{\text{CustID}=\text{CustID}'}(\text{Customer} \times \rho_{\text{CustID} \rightarrow \text{CustID}'}(\text{Account})))$$

where  $X = \{ \text{all attributes of Customer} \}$

$Y = \{ \text{all attributes of Account} \}$

# From SQL to relational algebra

**SELECT**  $\mapsto$  projection  $\pi$

**FROM**  $\mapsto$  Cartesian product  $\times$

**WHERE**  $\mapsto$  selection  $\sigma$

**SELECT**  $A_1, \dots, A_m$   
**FROM**  $T_1, \dots, T_n$   $\mapsto \pi_{A_1, \dots, A_m}(\sigma_{\langle \text{condition} \rangle}(T_1 \times \dots \times T_n))$   
**WHERE**  $\langle \text{condition} \rangle$

Common attributes in  $T_1, \dots, T_n$  must be renamed

## Set operations

### Union

R	A	B	$\cup$	S	A	B	=	R $\cup$ S	A	B
	a1	b1			a1	b1			a1	b1
	a2	b2			a3	b3			a2	b2
									a3	b3

### Intersection

R	A	B	$\cap$	S	A	B	=	R $\cap$ S	A	B
	a1	b1			a1	b1			a1	b1
	a2	b2			a3	b3				

### Difference

R	A	B	$-$	S	A	B	=	R $-$ S	A	B
	a1	b1			a1	b1			a2	b2
	a2	b2			a3	b3				

**The relations must have the same set of attributes**

## Union and renaming

R	Father	Child	S	Mother	Child
	George	Elizabeth		Elizabeth	Charles
	Philip	Charles		Elizabeth	Andrew
	Charles	William			

We want to find the relation **parent-child**

$\rho_{\text{Father} \rightarrow \text{Parent}}(\text{R}) \cup \rho_{\text{Mother} \rightarrow \text{Parent}}(\text{S})$	=	<table><tr><th>Parent</th><th>Child</th></tr><tr><td>George</td><td>Elizabeth</td></tr><tr><td>Philip</td><td>Charles</td></tr><tr><td>Charles</td><td>William</td></tr><tr><td>Elizabeth</td><td>Charles</td></tr><tr><td>Elizabeth</td><td>Andrew</td></tr></table>	Parent	Child	George	Elizabeth	Philip	Charles	Charles	William	Elizabeth	Charles	Elizabeth	Andrew
Parent	Child													
George	Elizabeth													
Philip	Charles													
Charles	William													
Elizabeth	Charles													
Elizabeth	Andrew													

## Full relational algebra

Primitive operations:  $\pi$  ,  $\sigma$  ,  $\times$  ,  $\rho$  ,  $\cup$  ,  $-$

Removing any of these results in a **loss of expressive power**

### Derived operations

$\bowtie$  can be expressed in terms of  $\pi$  ,  $\sigma$  ,  $\times$  ,  $\rho$

$\cap$  can be expressed in terms difference:

$$R \cap S \equiv R - (R - S)$$



## Other derived operations

Theta-join	$R \bowtie_{\theta} S \equiv \sigma_{\theta}(R \times S)$
Equijoin	$\bowtie_{\theta}$ where $\theta$ is a <b>conjunction of equalities</b>
Semijoin	$R \ltimes_{\theta} S \equiv \pi_X(R \bowtie_{\theta} S)$ where $X$ is the set of attributes of $R$
Antijoin	$R \bar{\ltimes}_{\theta} S \equiv R - (R \ltimes_{\theta} S)$

### Why use these operations?

- ▶ to write things more succinctly
- ▶ they can be optimized independently

## Division

$R$  over set of attributes  $X$

$S$  over set of attributes  $Y \subset X$

Let  $Z = X - Y$

$$\begin{aligned} R \div S &= \{ r \in \pi_Z(R) \mid \text{for every } s \in S, rs \in R \} \\ &= \{ r \in \pi_Z(R) \mid \{r\} \times S \subseteq R \} \\ &= \pi_Z(R) - \pi_Z(\pi_Z(R) \times S - R) \end{aligned}$$

Division: Example

Exams		DPT
Student	Course	Course
John	Databases	Databases
John	Networks	Programming
Mary	Programming	
Mary	Math	
Mary	Databases	

Exams ÷ DPT =

Student
Mary

$= \pi_{\text{Student}}(\text{Exams}) - \pi_{\text{Student}}(\pi_{\text{Student}}(\text{Exams}) \times \text{DPT} - \text{Exams})$