

# Transaction Management

Dr Paolo Guagliardo



THE UNIVERSITY of EDINBURGH  
**informatics**

Fall 2020 (v20.1.0)

## Transactions

**Transaction:** a sequence of operations on database objects


- ▶ All operations together form a **single logical unit**

### Example

Transfer £100 from account A to account B

1. Read balance from A into local buffer  $x$
2.  $x := x - 100$
3. Write new balance  $x$  to A
4. Read balance from B into local buffer  $y$
5.  $y := y + 100$
6. Write new balance  $y$  to B

# Life-cycle of a transaction



## Active

Normal execution state

## Partially Committed

Last statement executed

## Failed

Normal execution cannot proceed

## Aborted

Rolled-back

Previous database state restored

## Committed

Successful completion

Changes are permanent

## Schedules

**Schedule:** a sequence  $S$  of operations from a set of transactions s.t. the order of operations in each transaction is **the same** as in  $S$

A schedule is **serial** if all operations of each transaction are executed before or after all operations of another transaction

## Example

$T_1$  : op1, op2, op3

$T_2$  : op1, op2

### Concurrent schedule

	$T_1$	$T_2$
1		op1
2	op1	
3	op2	
4		op2
5	op3	

### Serial schedule

	$T_1$	$T_2$
1		op1
2		op2
3	op1	
4	op2	
5	op3	

# Concurrency

- ▶ Typically more than one transaction runs on a system
- ▶ Each transaction consists of many I/O and CPU operations
- ▶ We don't want to wait for a transaction to completely finish before executing another

## Concurrent execution

The operations of different transaction are **interleaved**

- ▶ increases throughput
- ▶ reduces response time

## The ACID properties

### Atomicity

Either all operations are carried out or none are

### Consistency

Successful execution of a transaction leaves the database in a coherent state

### Isolation

Each transaction is protected from the effects of other transactions executed concurrently

### Durability

On successful completion, changes persist

## Motivating example

$T_1$  : transfer £100 from account A to account B

$T_2$  : transfer 10% of account A to account B

$T_1$
1. $x := \text{read}(A)$
2. $x := x - 100$
3. $\text{write}(x, A)$
4. $y := \text{read}(B)$
5. $y := y + 100$
6. $\text{write}(y, B)$

$T_2$
1. $x := \text{read}(A)$
2. $y := 0.1 * x$
3. $x := x - y$
4. $\text{write}(x, A)$
5. $z := \text{read}(B)$
6. $z := z + y$
7. $\text{write}(z, B)$

$A + B$  should not change:

Money is not created and does not disappear

## Motivating example: Serial execution 1

	$T_1$	$T_2$	Database	
1	$x := \text{read}(A)$		$A = 1000$	$B = 1000$
2	$x := x - 100$		$A = 1000$	$B = 1000$
3	$\text{write}(x, A)$		$A = 900$	$B = 1000$
4	$y := \text{read}(B)$		$A = 900$	$B = 1000$
5	$y := y + 100$		$A = 900$	$B = 1000$
6	$\text{write}(y, B)$		$A = 900$	$B = 1100$
7		$x := \text{read}(A)$	$A = 900$	$B = 1100$
8		$y := 0.1 * x$	$A = 900$	$B = 1100$
9		$x := x - y$	$A = 900$	$B = 1100$
10		$\text{write}(x, A)$	$A = 810$	$B = 1100$
11		$z := \text{read}(B)$	$A = 810$	$B = 1100$
12		$z := z + y$	$A = 810$	$B = 1100$
13		$\text{write}(z, B)$	$A = 810$	$B = 1190$

## Motivating example: Serial execution 2

	$T_1$	$T_2$	Database	
1		$x := \text{read}(A)$	$A = 1000$	$B = 1000$
2		$y := 0.1 * x$	$A = 1000$	$B = 1000$
3		$x := x - y$	$A = 1000$	$B = 1000$
4		$\text{write}(x, A)$	$A = 900$	$B = 1000$
5		$z := \text{read}(B)$	$A = 900$	$B = 1000$
6		$z := z + y$	$A = 900$	$B = 1000$
7		$\text{write}(z, B)$	$A = 900$	$B = 1100$
8	$x := \text{read}(A)$		$A = 900$	$B = 1100$
9	$x := x - 100$		$A = 900$	$B = 1100$
10	$\text{write}(x, A)$		$A = 800$	$B = 1100$
11	$y := \text{read}(B)$		$A = 800$	$B = 1100$
12	$y := y + 100$		$A = 800$	$B = 1100$
13	$\text{write}(y, B)$		$A = 800$	$B = 1200$

## Motivating example: Concurrent execution 1

	$T_1$	$T_2$	Database	
1	$x := \text{read}(A)$		$A = 1000$	$B = 1000$
2	$x := x - 100$		$A = 1000$	$B = 1000$
3	$\text{write}(x, A)$		$A = 900$	$B = 1000$
4		$x := \text{read}(A)$	$A = 900$	$B = 1000$
5		$y := 0.1 * x$	$A = 900$	$B = 1000$
6		$x := x - y$	$A = 900$	$B = 1000$
7		$\text{write}(x, A)$	$A = 810$	$B = 1000$
8	$y := \text{read}(B)$		$A = 810$	$B = 1000$
9	$y := y + 100$		$A = 810$	$B = 1000$
10	$\text{write}(y, B)$		$A = 810$	$B = 1100$
11		$z := \text{read}(B)$	$A = 810$	$B = 1100$
12		$z := z + y$	$A = 810$	$B = 1100$
13		$\text{write}(z, B)$	$A = 810$	$B = 1190$

## Motivating example: Concurrent execution 2

	$T_1$	$T_2$	Database	
1	$x := \text{read}(A)$		$A = 1000$	$B = 1000$
2	$x := x - 100$		$A = 1000$	$B = 1000$
3		$x := \text{read}(A)$	$A = 1000$	$B = 1000$
4		$y := 0.1 * x$	$A = 1000$	$B = 1000$
5		$x := x - y$	$A = 1000$	$B = 1000$
6		$\text{write}(x, A)$	$A = 900$	$B = 1000$
7	$\text{write}(x, A)$		$A = 900$	$B = 1000$
8	$y := \text{read}(B)$		$A = 900$	$B = 1000$
9	$y := y + 100$		$A = 900$	$B = 1000$
10	$\text{write}(y, B)$		$A = 900$	$B = 1100$
11		$z := \text{read}(B)$	$A = 900$	$B = 1100$
12		$z := z + y$	$A = 900$	$B = 1100$
13		$\text{write}(z, B)$	$A = 900$	$B = 1200$

**We created £100 !!!**

## Transaction model

The only important operations in scheduling are **read** and **write**

$r(A)$  read data item  $A$

$w(A)$  write data item  $A$

Other operations do not affect the schedule

We represent transactions by a sequence of read/write operations

The transactions in the [motivating example](#) are represented as:

$T_1 : r(A), w(A), r(B), w(B)$

$T_2 : r(A), w(A), r(B), w(B)$

## Transaction model: Schedules

The schedules in the [motivating example](#) are represented as:

Schedule 1		Schedule 2	
$T_1$	$T_2$	$T_1$	$T_2$
$r(A)$		$r(A)$	
$w(A)$			$r(A)$
	$r(A)$		$w(A)$
	$w(A)$	$w(A)$	
$r(B)$		$r(B)$	
$w(B)$		$w(B)$	
	$r(B)$		$r(B)$
	$w(B)$		$w(B)$

Schedule 1 is **equivalent to a serial execution**, Schedule 2 is not

## Serializability

Two operations (from [different](#) transactions) are **conflicting** if

- ▶ they refer to the same data item, and
- ▶ at least one of them is a write

Two **consecutive** non-conflicting operations in a schedule can be [swapped](#)

A schedule is **conflict serializable** if it can be transformed into a serial schedule by a sequence of swap operations

# Precedence graph

Captures all potential conflicts between transactions in a schedule

- ▶ Each node is a transaction
- ▶ There is an edge from  $T_i$  to  $T_j$  (for  $T_i \neq T_j$ ) if an action of  $T_i$  **precedes** and **conflicts** with one of  $T_j$ 's actions

A schedule is **conflict serializable**

**if and only if**


its precedence **graph is acyclic**

An **equivalent serial schedule** is given by any **topological sort** over the precedence graph

## Precedence graph: Example

Schedule 1	
$T_1$	$T_2$
$r(A)$	
$w(A)$	
	$r(A)$
	$w(A)$
$r(B)$	
$w(B)$	
	$r(B)$
	$w(B)$

Precedence graph



Schedule 2	
$T_1$	$T_2$
$r(A)$	
	$r(A)$
	$w(A)$
$w(A)$	
$r(B)$	
$w(B)$	
	$r(B)$
	$w(B)$

Precedence graph





## Schedules with aborted transactions (1)

We assumed transactions commit successfully after the last operation

But **abort** and **commit** must be taken explicitly into account

	$T_1$	$T_2$
1	$r(A)$	
2	$w(A)$	
3		$r(A)$
4		$w(A)$
5		$r(B)$
6		$w(B)$
7	Abort	

- ▶  $T_2$  read uncommitted changes made by  $T_1$
- ▶ But  $T_2$  has not yet committed
- ▶ We can recover by aborting also  $T_2$   
(**cascading abort**)

## Schedules with aborted transactions (2)

	$T_1$	$T_2$
1	$r(A)$	
2	$w(A)$	
3		$r(A)$
4		$w(A)$
5		$r(B)$
6		$w(B)$
7		Commit
8	Abort	

- ▶  $T_2$  read uncommitted changes made by  $T_1$
- ▶ But  $T_2$  has already committed
- ▶ The schedule is **unrecoverable**

### Recoverable schedules without cascading aborts

Transactions commit only after, and if,  
all transactions whose changes they read commit

# Lock-based concurrency control

## Lock

- ▶ Bookkeeping object associated with a data item
- ▶ Tells whether the data item is available for read and/or write
- ▶ **Owner**: Transaction currently operating on the data item

**Shared lock** Data item is available for read to owner  
Can be acquired by more than one transaction

**Exclusive lock** Data item is available for read/write to owner  
Cannot be acquired by other transactions

Two locks on the same data item are **conflicting**  
if one of them is exclusive

## Transaction model with locks

Operations:

$s(A)$  **shared lock** on  $A$  is acquired

$x(A)$  **exclusive lock** on  $A$  is acquired

$u(A)$  lock on  $A$  is released

**Abort** transaction aborts

**Commit** transaction commits

In a schedule:

- ▶ A transaction cannot acquire a lock on  $A$   
before all exclusive locks on  $A$  have been released
- ▶ A transaction cannot acquire an exclusive lock on  $A$   
before all locks on  $A$  have been released

## Examples of schedules with locking

Schedule 1		Schedule 2	
$T_1$	$T_2$	$T_1$	$T_2$
$x(A)$		$s(A)$	
$u(A)$			$s(A)$
	$x(A)$		$u(A)$
	$u(A)$	$u(A)$	
$x(B)$			$x(A)$
$u(B)$			$u(A)$
Commit		$x(A)$	
	$x(B)$	$x(B)$	
	$u(B)$	$u(B)$	
	Commit		$x(B)$
			$u(B)$
			Commit
		$u(A)$	
		Commit	

## Two-Phase Locking (2PL)

1. Before reading/writing a data item  
a transaction must acquire a shared/exclusive lock on it
2. A transaction cannot request additional locks  
once it releases **any** lock

Each transaction has

**Growing phase** when locks are acquired

**Shrinking phase** when locks are released

Every completed schedule of **committed** transactions  
that follow the 2PL protocol is conflict serializable

## 2PL and aborted transactions

	$T_1$	$T_2$
1	$x(A)$	
2	$u(A)$	
3		$x(A)$
4		$x(B)$
5		$u(A)$
6		$u(B)$
7		Commit
8	Abort	

- ▶  $T_1$  and  $T_2$  follow 2PL
- ▶ But  $T_1$  cannot be undone
- ▶ The schedule is **unrecoverable**

## Strict 2PL

1. Before reading/writing a data item  
a transaction must acquire a shared/exclusive lock on it
2. **All locks held by a transaction are released  
when the transaction is completed** (aborts or commits)

Ensures that

- ▶ The schedule is always **recoverable**
- ▶ All aborted transactions can be rolled back  
without cascading aborts
- ▶ The schedule consisting of the committed transactions  
is **conflict serializable**

## Deadlocks

A transaction requesting a lock must wait until all conflicting locks are released

We may get a cycle of “waits”

	$T_1$	$T_2$	$T_3$
1	s(A)		
2		x(B)	
3	req s(B)		
4			s(C)
5		req x(C)	
6			req x(A)

$T_1$  waits for  $T_2$  ,    $T_2$  waits for  $T_3$  ,    $T_3$  waits for  $T_1$

## Deadlock prevention

Each transaction is assigned a **priority** using a **timestamp**:  
The older a transaction is, the higher priority it has

Suppose  $T_i$  requests a lock and  $T_j$  holds a conflicting lock

Two policies to prevent deadlocks:

**Wait-die:**  $T_i$  waits if it has higher priority, otherwise aborted

**Wound-wait:**  $T_j$  aborted if  $T_i$  has higher priority, otherwise  $T_i$  waits

In both schemes, the higher priority transaction is never aborted

**Starvation:** a transaction keeps being aborted  
because it never has sufficiently high priority

**Solution:** restart aborted transactions with their initial timestamp

# Deadlock detection

## Waits-for graph

- ▶ Nodes are active transactions
- ▶ There is an edge from  $T_i$  to  $T_j$  (with  $T_i \neq T_j$ ) if  $T_i$  waits for  $T_j$  to release a (conflicting) lock

Each cycle represents a deadlock

## Recovering from deadlocks

Choose a minimal set of transactions such that rolling them back will make the waits-for graph acyclic

# Crash recovery

## The log (a.k.a. **trail** or **journal**)

Records every action executed on the database

Each log record has a unique ID called **log sequence number (LSN)**

Fields in a log record:

**LSN** ID of the record

**prevLSN** LSN of previous log record

**transID** ID of the transaction

**type** of action recorded

**before** value before the change

**after** value after the change

The state of the database is periodically recorded as a **checkpoint**

# ARIES

Recovery algorithm used in major DBMSs

Works in three phases

## 1. Analysis

- ▶ identify changes that have not been written to disk
- ▶ identify active transactions at the time of crash

## 2. Redo

- ▶ repeat all actions starting from latest checkpoint
- ▶ restore the database to the state at the time of crash

## 3. Undo

- ▶ undo actions of transactions that did not commit
- ▶ the database reflects only actions of committed transactions

# Principles behind ARIES

## Write-Ahead Logging

Before writing a change to disk, a corresponding log record must be inserted and the log forced to stable storable

## Repeating history during Redo

Actions before the crash are retraced to bring the database to the state it was when the system crashed

## Logging changes during Undo

Changes made while undoing transactions are also logged (protection from further crashes)