



THE UNIVERSITY of EDINBURGH
informatics

Operating Systems (INFR10079) 2020/2021 Semester 2

Structure (Operating System Structure)

abarbala@inf.ed.ac.uk


Chapter 2.4, 2.7, 2.8

Overview

- *Architecture impact*
- *Application-Operating System interaction*
- **Operating System structure**

OS Structure

- OS mediates access and abstracts away ugliness
- OS sits between **applications** and the **hardware**
 - Applications (**App**) request services
 - **Explicitly** via syscalls
 - **Implicitly** via exceptions
 - Devices (**Device**) request attention via interrupts



Operating System Design and Implementation


- Design and Implementation of OS not “solvable”
 - but some approaches have proven successful
- Internal structure of different Operating Systems can vary widely
- Start the design by defining goals and specifications
 - **User** goals: convenient to use, easy to learn, reliable, safe, and fast
 - **System** goals: easy to design, implement, and maintain, as well as flexible, reliable, error-free, and efficient
- Affected by choice of hardware, type of system

Operating System Design and Implementation

- Important **principles to separate**
- **Policy:** *What* will be done? (Algorithm)
- **Mechanism:** *How* to do it?
- Separation allows maximum flexibility
 - Policies are likely to change across places or over time
 - A general mechanism can support a wide range of policies
- Microkernel OSes are based on such principle
 - A core kernel implements the mechanisms
 - Policies are implemented outside the core kernel
 - Easily modifiable

Major OS Services

- processes
- memory
- I/O
- secondary storage
- file systems
- protection
- networking
- shells (i.e., command interpreter)
- GUI
- etc.



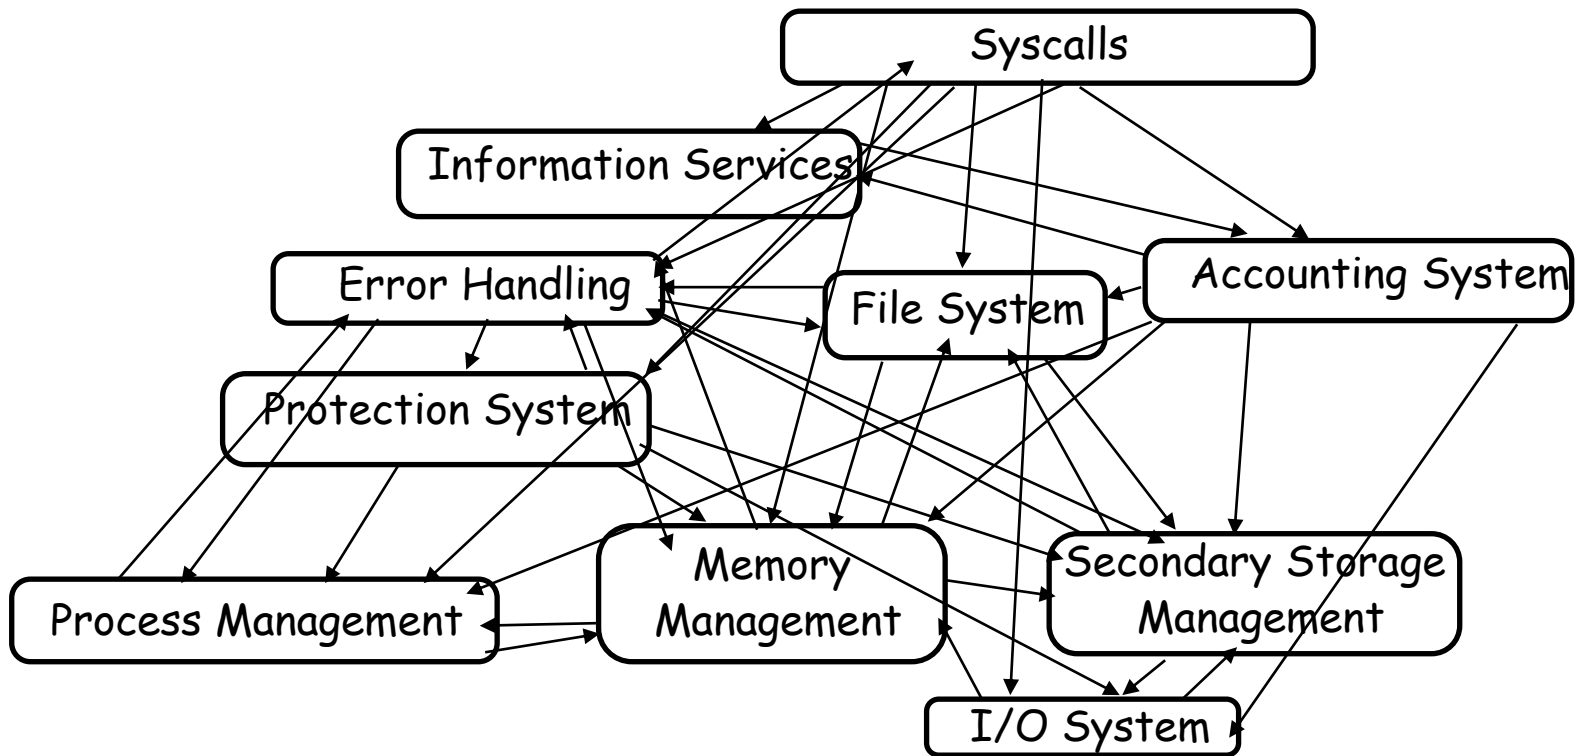
```
#!/bin/bash
```

```
~root: env X="()" { : } ; echo shellshock" /bin/sh -c "echo completed"  
> shellshock  
> completed
```

Systems programs – outside the kernel

OS Structure #1

- It's not always clear how to stitch OS **services** together




OS Structure #2


- Major issues
 - how do we organize all these?
 - what are all of the code modules, and where do they exist?
 - how do they cooperate?
- Massive software engineering and design problem
 - design a large, complex program that
 - performs well
 - is reliable
 - is extensible
 - is backwards compatible
 - etc.

Monolithic OS Design #1

- Likely the earliest OS organization
- **UNIX** was built as **monolithic**
 - **Linux** is built as monolithic



Monolithic Example: Linux



Monolithic OS Design #2

- Major **advantage**
 - **cost** of subsystems interactions **is low** (procedure call)
- Disadvantages
 - hard to understand
 - hard to modify
 - unreliable (no isolation between system modules)
 - hard to maintain
- What is the alternative?
 - find a way to organize OS subsystems to simplify its design and implementation


Layered OS Design

- The traditional approach is layering
 - implement OS as a set of layers
 - each layer presents an enhanced 'virtual machine' to the layer above
- The first description of this approach was Dijkstra's THE system
 - Layer 5: **Job Managers**
 - Execute users' programs
 - Layer 4: **Device Managers**
 - Handle devices and provide buffering
 - Layer 3: **Console Manager**
 - Implements virtual consoles
 - Layer 2: **Page Manager**
 - Implements virtual memories for each process
 - Layer 1: **Kernel**
 - Implements a virtual processor for each process
 - Layer 0: **Hardware**
- Each layer can be tested and verified independently



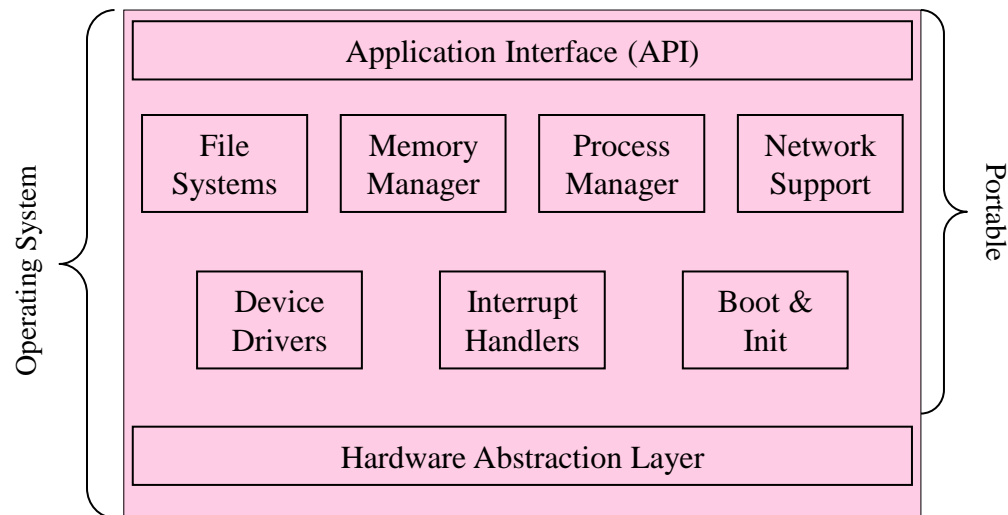
Problems with layering

- **Imposes** hierarchical structure
 - but real systems are more complex
 - File system requires virtual memory services
 - Virtual memory would like to use files for its backing store
 - **strict layering isn't flexible enough**
- Poor performance
 - each layer crossing has **overhead** associated with it
- Disjunction between model and reality
 - systems modeled as layers, but not really built that way



Hardware Abstraction Layer

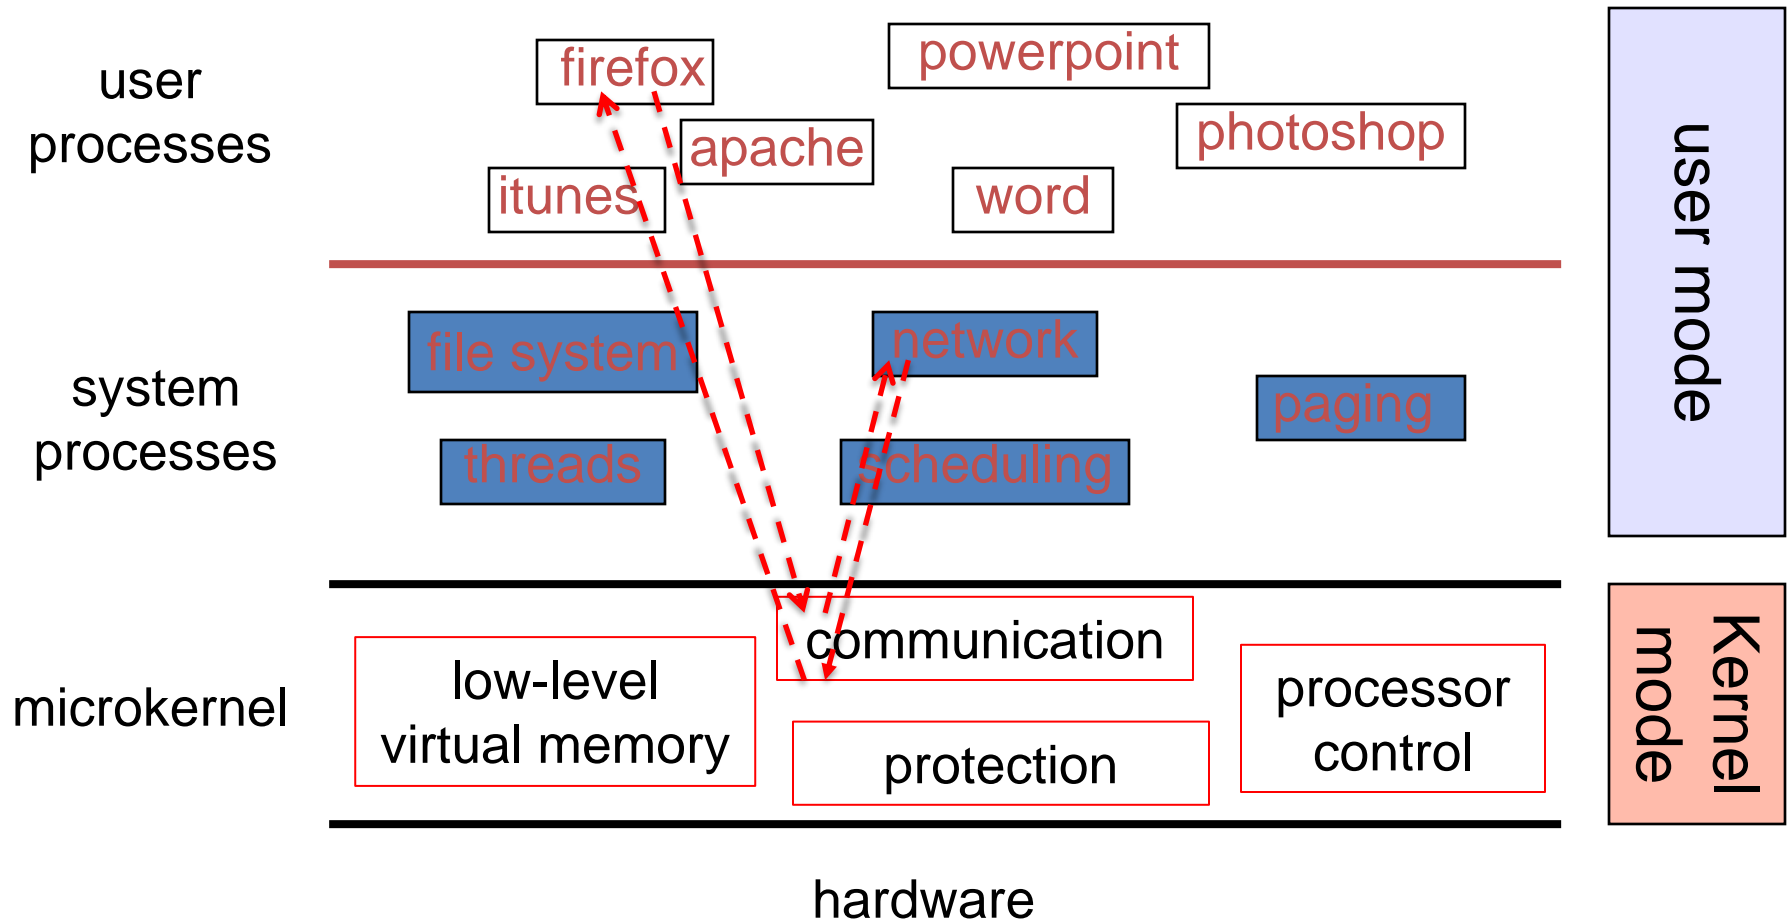
- An **example** of layering in modern operating systems
 - Windows, etc.
- Goal: separates hardware-specific routines from the **core kernel** of the OS
 - Provides portability
 - Improves readability



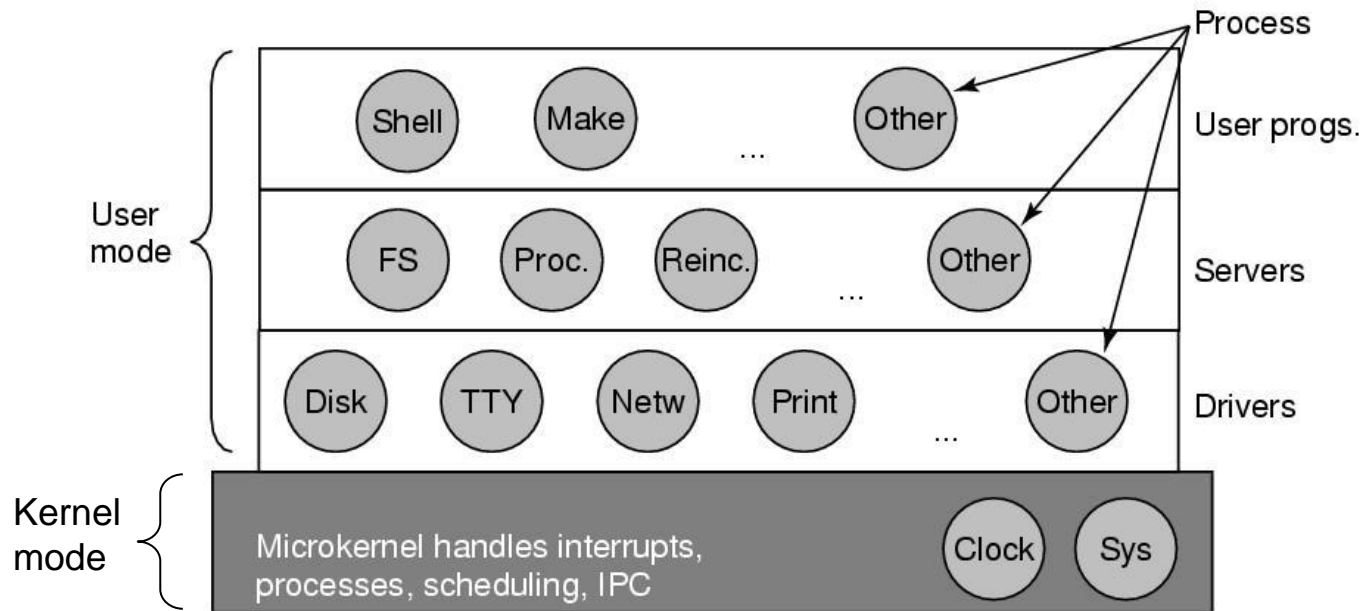
Microkernel OS Design

- Popular in the late 80's, early 90's
 - **recent** resurgence of popularity
- Goal
 - **minimize** what goes in kernel
 - organize rest of OS as user-level processes (services)
- This results in
 - better **reliability** (isolation between components)
 - ease of **extension and customization**
 - **poor performance** (user/kernel boundary crossings)
- First microkernel system was Hydra (CMU, 1970)
 - Follow-ons: Mach (CMU), Chorus (French UNIX-like OS), MINIX (UNIX-like OS from Amsterdam)


Microkernel Structure Illustrated



Microkernel Example: MINIX




Loadable Kernel Modules



- **Core services** in the kernel, **others** dynamically loaded
- Common in modern implementations
 - **Monolithic**: load the code in kernel space (Solaris, Linux, etc.)
 - **Microkernel**: load the code in user space (any)
- Advantages
 - **Convenient**: no need for rebooting for newly added modules
 - **Efficient**: no need for message passing unlike microkernel
 - **Flexible**: any module can call any other module unlike layered model

Hybrid OS Design

- Many **different** approaches
 - Key idea: exploit the benefits of monolithic and microkernel designs
 - Windows, Xnu/Darwin, DragonFly BSD, ...
- Extensibility via kernel modules



Picture Copyright of Wikipedia

Summary

- Fundamental distinction between user and privileged modes supported by most hardware
- OS design has been an evolutionary process of trial and error
- Successful OS designs have run the spectrum from monolithic, to layered, to micro kernels
- The role and design of an OS are still evolving
- There is no “ideal” OS structure