



THE UNIVERSITY of EDINBURGH  
**informatics**

# **Operating Systems (INFR10079) 2020/2021 Semester 2**

## **Processes (Interprocess Communication)**


[abarbala@inf.ed.ac.uk](mailto:abarbala@inf.ed.ac.uk)

Chapter 3.4, 3.5, 3.6, 3.7, 3.8, excluding: 3.7.2, 3.7.3, 3.8.2.1  
4.6.2, 20.9.1, C3.3

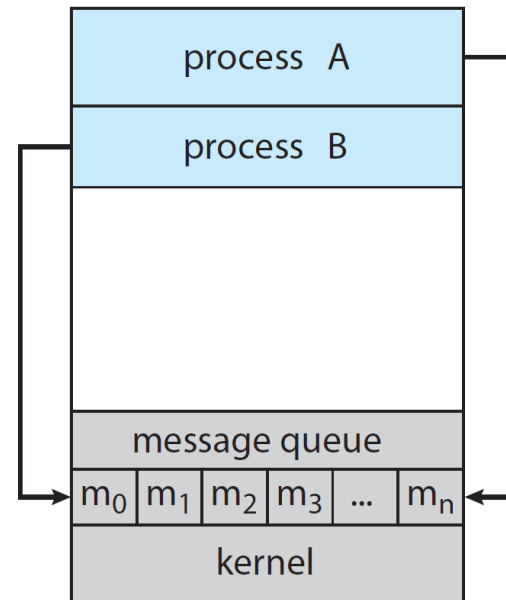
# Interprocess Communication

- **Independent** processes
- **Cooperating** processes
  - Information **sharing**
    - More applications interested in the same information
  - Computation **speedup**
    - To exploit parallel hardware
  - **Modularity**
    - Reusability of components
- Require **interprocess communication** (IPC) mechanism to send and receive data
  - Shared memory
  - Message passing

# Interprocess Communication



Shared memory



Message passing

# Shared Memory

- Allow processes to communicate and synchronize
  - **Sharing** part of address space
  - **OS doesn't mediate** communication (no overhead)
    - Usually, **OS prevents** processes to access each other memory
    - Processes should agree (have permission) to void this restriction
- Data
  - Format **decided by application**
  - Direct access (not mediated by the OS) – **very fast**
  - Application programmer fully manages the data transfer – **not trivial**
- Possible use cases
  - Passing of large (single) objects (C&P, image, table, etc.)
  - Notification variable

# Example of Shared Memory Code (POSIX)

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <fcntl.h>
#include <sys/shm.h>
#include <sys/stat.h>
#include <sys/mman.h>

int main() {
    const int SIZE = 4096; /* the size (bytes) of shared memory object */
    const char *name = "OS"; /* name of the shared memory object */
    const char *message 0 = "Hello"; /* written to shared memory */
    const char *message 1 = "World!"; /* written to shared memory */
    int fd; /* shared memory file descriptor */
    char *ptr; /* pointer to shared memory object */

    /* create the shared memory object */
    fd = shm_open(name, O_CREAT | O_RDWR, 0666);
    /* configure the size of the shared memory object */
    ftruncate(fd, SIZE);
    /* memory map the shared memory object */
    ptr = (char *)
        mmap(0, SIZE, PROT_READ | PROT_WRITE, MAP_SHARED, fd, 0);
    /* write to the shared memory object */
    sprintf(ptr, "%s", message 0);
    ptr += strlen(message 0);
    sprintf(ptr, "%s", message 1);
    ptr += strlen(message 1);

    return 0;
}
```

```
#include <stdio.h>
#include <stdlib.h>
#include <fcntl.h>
#include <sys/shm.h>
#include <sys/stat.h>
#include <sys/mman.h>

int main() {
    const int SIZE = 4096; /* the size (bytes) of shared memory object */
    const char *name = "OS"; /* name of the shared memory object */

    int fd; /* shared memory file descriptor */
    char *ptr; /* pointer to shared memory object */

    /* open the shared memory object */
    fd = shm_open(name, O_RDONLY, 0666);
    /* memory map the shared memory object */
    ptr = (char *)
        mmap(0, SIZE, PROT_READ | PROT_WRITE, MAP_SHARED, fd, 0);
    /* read from the shared memory object */
    printf("%s", (char *)ptr);
    /* remove the shared memory object */
    shm_unlink(name);

    return 0;
}
```

# Message Passing

- Allow processes to communicate and synchronize
  - **Without sharing** part of address space
  - **OS mediates** communication (may introduce overhead)
- Processes on the **same** machine and **among different** inter-networked machines
  - Not possible with shared memory
- Message-passing facility provides at least two operations
  - **send(message)**
  - **receive(message)**
- Communication link
  - Several **implementation tradeoffs**, e.g., messages size
    - Fixed
    - Variable

# Message Passing – Naming

- Communicating processes must **refer to each other**
- **Direct** communication
  - **Symmetric:** Explicit name of sender and receiver
    - `send(P, message)` – send message to process P
    - `receive(Q, message)` – receive message from process Q
  - **Asymmetric:** Explicit at least one end
    - `send(p, message)` – send to p
    - `receive(id, message)` – receive from any process, sender saved in id
- **Indirect** communication
  - No need to know/explicit in advance sender and/or receiver
  - Mailboxes (e.g., POSIX Mailbox)
    - `send(A, message)` – into mailbox A
    - `receive(A, message)` – from mailbox A
  - A mailbox can be accessed by more than two processes
    - Between processes multiple mailboxes may exist



Naming is a  
problem  
also in  
shared  
memory

# Message Passing – Synchronization

- Different **design options** to implement `send()` / `receive()`
  - **Blocking** or **synchronous**
  - **Nonblocking** or **asynchronous**
- Different **combinations** may be offered
  - Blocking send
  - Nonblocking send
  - Blocking receive
  - Nonblocking receive
- **Rendezvous**
  - When both `send()` and `receive()` **are blocking**
  - Simple solution to the consumer-producer problem\*\*\*




# Message Passing – Buffering

- Message exchanged reside in **temporary buffers/queues**
  - **Zero** capacity (no buffering)
    - No message waiting in it
    - Sender must block until the recipient receives the message
  - **Bounded** capacity
    - Finite length ***n***; at most n messages can reside in it
    - If the queue is not full when a new message is sent, that is placed in the queue and the sender can continue execution
    - If the link is full, the sender has to wait
  - **Unbounded** capacity
    - Infinite queue
    - The sender never blocks

# Example of message passing: Pipes

- A pipe acts as a **conduit** allowing two processes to communicate
  - One way data flow
- **Anonymous** (ordinary)
  - Between parent and child
- **Named**
  - between any pair of processes)
  - In UNIX, FIFO



**Figure 3.20** File descriptors for an ordinary pipe.

# Example of Pipe Code (POSIX)

```
#include <sys/types.h>
#include <stdio.h>
#include <string.h>
#include <unistd.h>
#define READ_END 0
#define WRITE_END 1
int main(void) {
    char write_msg[256] = "Greetings";
    char read_msg[256];
    int fd[2];
    pid_t pid;

    if (pipe(fd) == -1) { /* create the pipe */
        fprintf(stderr, "Pipe failed");
        return 1;
    }
    pid = fork(); /* fork a child process */
    if (pid < 0) { /* error occurred */
        fprintf(stderr, "Fork Failed");
        return 1;
    }
    if (pid > 0) { /* parent process */
        close(fd[READ_END]); /* close the unused end of the pipe */
        write(fd[WRITE_END], write_msg, strlen(write_msg)+1); /* write to the pipe */
        close(fd[WRITE_END]); /* close the write end of the pipe */
    }
    else { /* child process */
        close(fd[WRITE_END]); /* close the unused end of the pipe */
        read(fd[READ_END], read_msg, 256); /* read from the pipe */
        printf("read %s", read_msg);
        close(fd[READ_END]); /* close the read end of the pipe */
    }
    return 0;
}
```

# Client-Server Communication

- **Sockets** Abstraction
  - Endpoint for communication
  - Identified by an **IP address** concatenated with a **port number**
  - Servers implementing specific services (SSH, FTP, and HTTP) listen to well-known ports
    - an SSH server listens to port 22; an FTP server listens to port 21; and a web, or HTTP, server listens to port 80
- **Remote Procedure Call (RPC)**
  - Abstract the **procedure-call mechanism**
  - For use between systems with network connections
  - **Similar** in many respects **to the IPC**
    - Uses message-based communication to provide remote service

# Signals

- **OS mechanism to notify a process** (one way)
  - Doesn't carry information
- From **the OS**
  - Can be thought as a **software-generated interrupt/exception**
    - Synchronous – e.g., division by zero error
    - Asynchronous – e.g., timer/alarm
  - **Syscall: process -> OS; Signal: OS -> process**
- From **other processes** (including the process itself)
  - Only notification, no data **communication method**
    - Management – e.g., kill a process
    - Synchronization – e.g., POSIX RT Signals
    - etc.

# Handling Signals (Linux Example)

- **Signal handler**
  - Code to process a signal
  - No return value (void)
- Handler must be registered
  - Or OS default action
- Send a signal
  - `kill()`
  - `raise()`

[https://www.cs.auckland.ac.nz/references/unix/digital/APS33DTE/DOCU\\_006.HTM#realtime-handler-sec](https://www.cs.auckland.ac.nz/references/unix/digital/APS33DTE/DOCU_006.HTM#realtime-handler-sec)

```
#include <unistd.h>
#include <signal.h>
#include <stdio.h>
#include <sys/types.h>
#include <sys/wait.h>
#define SIG_STOP_CHILD SIGUSR1

main() {
    pid_t pid;
    sigset_t newmask, oldmask;

    if ((pid = fork()) == 0) { /* Child */
        struct sigaction action;
        void catchit();
        sigemptyset(&newmask);
        sigaddset(&newmask, SIG_STOP_CHILD);
        sigprocmask(SIG_BLOCK, &newmask, &oldmask);
        action.sa_flags = 0;
        action.sa_handler = catchit;
        if (sigaction(SIG_STOP_CHILD, &action, NULL) == -1) {
            perror("sigusr: sigaction");
            _exit(1);
        }
        sigsuspend(&oldmask);
    }
    else { /* Parent */
        int stat;
        sleep(10);
        kill(pid, SIG_STOP_CHILD);
        pid = wait(&stat);
        printf("Child exit status = %d\n", WEXITSTATUS(stat));
        _exit(0);
    }
}

void catchit(int signo) { /* Signal Handler */
    printf("Signal %d received from parent\n", signo);
    _exit(0);
}
```